University Comparison Framework

So basically I have a question to compare four universities

CSUB, CSLB, UCLA, UCB

(Phase S — State and clarify)

Hi. I'm your thinking partner for Assignment 4 — the ANOVA and TukeyHSD problem with GPA data from four universities. A few ground rules so this works:

- I coach; I don't answer. You'll do the computing and R work.
- We follow six phases (SMARTS). I'll tell you which one we're in.
- I need a **FRAME-style opening** from you before we go further:
  Facts, **R**esearch objective, **A**ssumptions, **M**ethod, **E**xact sticking point.
- This is for academic integrity — the goal is that *you* understand the ANOVA and Tukey steps.

You've said you want to compare CSUB, CSLB, UCLA, and UCB — good start, but that's not a full FRAME yet.

Start with:

- **Facts**: what data do you have (groups, sample sizes, variable)?
- **Research objective**: what exactly are you testing?
- **Assumptions**: what conditions ANOVA needs
- **Method**: what test(s) you plan to use
- **Exact sticking point**: where you're currently stuck

Write your FRAME, and I'll check it.

hmm okay
CSUB, CSLB, UCLA, UCB

and

```
CSUB <- c(3.0, 2.9, 3.1, 3.2, 3.0, 3.1, 2.8, 3.0, 3.2)
CSLB <- c(3.4, 3.5, 3.3, 3.6, 3.4, 3.3, 3.5, 3.6)
UCLA <- c(3.8, 3.9, 3.7, 3.9, 4.0, 3.8, 3.7, 3.9, 3.8)
UCB  <- c(3.6, 3.7, 3.5, 3.6, 3.7, 3.5, 3.6, 3.8, 3.7)
```

this is what I was trying to enter r

basically want to see which schools gpa differ and questions asks something about a TukeyHSD too

(Phase S — State and clarify)

You've got the raw data — good. But this still isn't a full FRAME yet. Right now you've mostly given me **Facts (partial)** and a vague goal.

Let's tighten it up. I'm going to push you a bit here.

Fix your FRAME:

1. Facts (complete this)
You need to explicitly state:

- What is the **response variable**?
- How many **groups (k)**?
- What are the **sample sizes** for each school?

👉 Try writing that in one clean sentence.

---

2. Research objective (too vague right now)
You said: *"see which schools GPA differ"* — that's only half.

You need **two parts**:

- the **overall test** (ANOVA)
- the **follow-up comparisons** (Tukey)

👉 Write one sentence starting like:

> "Test whether … and if so …"

---

3. Assumptions (missing completely)
List the 3 ANOVA assumptions. Don't explain deeply yet — just name them.

---

4. Method (partially there)
You mentioned TukeyHSD, good. But I want:

- What is the **main test called**?
- When do you use Tukey?

---

5. Exact sticking point (this is key)
Right now you said:

> "trying to enter r"

That's not specific enough.

👉 Are you stuck on:

- combining the data?
- running `aov()`?
- interpreting output?
- understanding ANOVA conceptually?

Your turn

Rewrite your full FRAME with all 5 parts. Keep it short but precise.

Once that's solid, we'll move to **Phase M (plan)** and then I'll help you with the R step *without just handing it to you*.

so basically run f test

then if it says atleast one is diff check something else?

(Phase S — State and clarify)

You're **close**, but what you said is still too informal for the assignment. Let's sharpen it.

What you wrote:

> "run f test then if at least one is diff check something else"

That's the *idea*, but I need you to state it in proper statistical terms.

I want you to rewrite just the Research objective + Method like this:

👉 Fill in the blanks:

Research objective:
"Test whether _____, and if not, determine _____."

Method:
"We will use a _____, and if we reject the null, we will use _____ to identify which pairs differ."

One hint (only one, per the rules):

- The F-test only tells you: **"at least one mean is different"**
- It does **NOT** tell you *which ones*

Fix those two pieces. Don't move on yet.

Once that's correct, we'll lock in your FRAME and go to **Phase M (plan)**.

at least one mean is diff and then figure out which ones are actually diff?

(Phase S — State and clarify)

That's the right *idea*, but tighten the wording so it sounds like a proper statistical objective.

Right now:

> "at least one mean is diff and then figure out which ones are actually diff"

Let's make it precise.

Refine it like this:

👉 Try again:

Research objective:
"Test whether the _____, and if not, determine _____."

(Use phrases like *population mean GPA* and *which pairs of universities differ*)

Small push:

- Say **"all four population means are equal"** instead of "means"
- Say **"which pairs of universities differ"** instead of "which ones"

Now also add the **Method** right under it:

👉 Fill this in:
"We will use a _____, and if we reject the null hypothesis, we will use _____."

Write both cleanly. Once that's solid, we move on.

all four population means are equal and which pairs of universities differ"

we will first do F test then if rejects do tukey HSD?

(Phase S — State and clarify)

Good — now we're getting somewhere. The structure is right, just clean up the wording so it's fully correct.

Fix your statements:

Research objective (final form):
You need one complete sentence:

👉 Try this:

> "Test whether the four population mean GPAs are equal across the universities, and if not, determine which pairs of universities differ."

Method (tighten terminology):
What you wrote:

> "we will first do F test then if rejects do tukey HSD"

Close, but name things properly:

👉 Rewrite as:

> "We will use a one-way ANOVA F-test, and if we reject the null hypothesis, we will use Tukey's HSD test to identify which pairs differ."

One correction (important):

Don't just say "F test" — say **"one-way ANOVA F-test"**. That matters.

Now we're still missing:

You haven't done:

- **Assumptions**
- **Exact sticking point**

Your turn

Add:

Assumptions: (list the 3)
Exact sticking point: (be specific — R? formulas? interpretation?)

Once that's complete, we'll move to **Phase M (plan)** and start building the ANOVA step.

I want to run the F test now.

```
> aov(gpa_data ~ gpa, data = gpa_data)
Error in model.frame.default(formula = gpa_data ~ gpa, data = gpa_data,  :
  invalid type (list) for variable 'gpa_data'
>
```

got an error

(Phase A — Ask and gather)

The error is from the **formula**.

You wrote:

```R
aov(gpa_data ~ gpa, data = gpa_data)
```

That tells R to use the whole data frame `gpa_data` as the response variable, which is why it says `invalid type (list)`.

For one-way ANOVA, the formula shape is:

```R
response ~ group
```

So ask yourself:

- What is the **response** here?
- What is the **grouping variable** here?

From your data frame, those should be:

- response = `gpa`
- group = `university`

So your formula should have the form:

```R
aov(gpa ~ university, data = gpa_data)
```

Now you do this part:

1. Run that corrected `aov(...)`
2. Then run `summary(...)` on the model
3. Paste the output here

Do **not** interpret it yet. First just show me the ANOVA table.

I got this but has not info I need

```
aov(gpa ~ university, data = gpa_data)
Call:
   aov(formula = gpa ~ university, data = gpa_data)

Terms:
                university Residuals
Sum of Squares    3.135429  0.400000
Deg. of Freedom          3        31

Residual standard error: 0.1135924
Estimated effects may be unbalanced
>
```

like p value etc

(Phase R — Reason step by step)

Good — you ran the model correctly. The issue now is just **how you're extracting the ANOVA table**.

What you printed:

```R
aov(gpa ~ university, data = gpa_data)
```

That only shows the **structure**, not the full test results.

Key idea:

The **F-statistic and p-value** come from the **ANOVA table**, which you get using:

```R
summary(model)
```

But only if you *store the model first*.

Fix your steps:

👉 Step 1: Save the model

```R
model <- aov(gpa ~ university, data = gpa_data)
```

👉 Step 2: Then run

```R
summary(model)
```

Your turn

Run those two lines and paste the output.

Don't interpret it yet — we'll walk through:

- where the F value is
- where the p-value is
- what each row means

step by step.

```
> summary(model)
Error: object 'model' not found

error here
```

(Phase R — Reason step by step)

That error just means the model was never saved — probably because you didn't run the line in order.

R works line-by-line. If `model` doesn't exist, `summary(model)` can't work.

Let's fix it cleanly:

👉 Run **these two lines in order** (one at a time):

```R
model <- aov(gpa ~ university, data = gpa_data)
```

then:

```R
summary(model)
```

Quick check:

After the first line, you shouldn't see much output — that's normal.
The **second line** is what gives you:

- F value
- p-value

Your turn

Run both lines again carefully and paste the output.

Don't skip steps — ANOVA is very procedural.

got it pvalue is soooooo tiny we go to tukey? explainf irst

(Phase T — Think critically and iterate)

Good — now slow down. You said:

> "p-value is so tiny → we go to Tukey"

That's *procedurally* correct, but I want you to understand **why** before we move on.

First question (don't skip this):

👉 What does a **very small p-value** in the ANOVA F-test actually tell you about the null hypothesis?

Be specific:

- Does it say *all means are different*?
- Or something weaker?

Second question (this is the key idea of Tukey):

Suppose ANOVA tells you:

> "at least one mean is different"

👉 Why can't we just:

- look at the sample means, or
- run a bunch of regular t-tests?

What problem happens if you compare all pairs without adjustment?

(Hint: it has to do with making **multiple comparisons**.)

Answer those two in your own words. Once you've got that, we'll move to **running TukeyHSD correctly in R**.

so now we know at least one is different we want to confirm which one right?

(Phase T — Think critically and iterate)

Yes — that's the idea, but tighten it:

> ANOVA tells us **at least one population mean is different**, but it does **not tell us which ones**.

So your statement becomes:

> "Now we want to determine **which specific pairs of universities differ**."

But here's the deeper part (this matters):

👉 Why don't we just run a bunch of regular t-tests?

If you compare all 4 schools:

- How many pairwise comparisons are there?

Try to count them:
CSUB vs CSLB, CSUB vs UCLA, … etc.

👉 How many total?

Key concept I want you to say:

Running many tests increases the chance of making a **Type I error** (false positive).

That's exactly why we use **TukeyHSD** — it controls that.

Your turn

1. How many pairwise comparisons are there with 4 groups?
2. What problem happens if we run that many t-tests at α = 0.05?

Answer both, then we'll run **TukeyHSD in R**.

> 6 comparisons?

(Phase T — Think critically and iterate)

Yes — **6 comparisons** is correct.

Now finish the reasoning:

👉 If you run 6 separate t-tests at α = 0.05…

- Each test has a 5% chance of a false positive
- Across 6 tests, what happens to the **overall chance** of making *at least one* false positive?

Does it:

- stay at 0.05, or
- get larger?

I want you to say this idea clearly:

Running multiple tests **increases the probability of a Type I error** (this is called the **family-wise error rate**).

One more step:

👉 So why is **TukeyHSD better** than just doing 6 t-tests?

Answer that in one sentence. Then we move to actually running it in R.

yes

(Phase T — Think critically and iterate)

Not just "yes." Say the idea.

Write one sentence like this in your own words:

> "If we run 6 separate t-tests, the overall chance of a false positive gets larger than 0.05, so we use TukeyHSD to control the family-wise error rate while comparing all pairs."

Once you type that, the next R step is:

R
```
TukeyHSD(model)
```

Run it after you state the reason, then paste the output here.

```
             diff        lwr       upr     p adj
CSUB-CSLB -0.4166667 -0.56647250 -0.2668608 0.0000001
UCB-CSLB   0.1833333  0.03352750  0.3331392 0.0117123
UCLA-CSLB  0.3833333  0.23352750  0.5331392 0.0000005
UCB-CSUB   0.6000000  0.45466699  0.7453330 0.0000000
UCLA-CSUB  0.8000000  0.65466699  0.9453330 0.0000000
UCLA-UCB   0.2000000  0.05466699  0.3453330 0.0040144
```

okay so except UCB-CSLB and UCLA-UCB everything else is dig different?

(Phase R — Reason step by step)

Careful — you jumped too fast and made a mistake.

You said:

> "except UCB-CSLB and UCLA-UCB everything else is different"

That's **not correct**.

Look at the rule for Tukey:

A pair is **significantly different if**:

- the **p adj < 0.05**, OR
- the **confidence interval does NOT include 0**

Now check your two "exceptions":

UCB – CSLB

- p adj = 0.0117 → **less than 0.05**
- CI = (0.0335, 0.3331) → does NOT include 0
  👉 **This IS significant**

UCLA – UCB

- p adj = 0.0040 → **less than 0.05**
- CI = (0.0547, 0.3453) → does NOT include 0
  👉 **This IS significant**

So what's the correct conclusion?

👉 Look at all 6 rows:

- Are any p-values **greater than 0.05**?
- Do any confidence intervals include **0**?

Your turn

Answer this clearly:

👉 "How many of the 6 pairs are statistically significantly different?"

Then we'll move to **final interpretation (Phase S)**.

so all are different?

(Phase S — Summarize and verify)

Yes — **all 6 pairs are statistically significantly different**.

You verified it correctly:

- Every **p adj < 0.05**
- Every **confidence interval does NOT include 0**

Now don't stop at "all are different" — finish the interpretation:

👉 What is the **ordering of the mean GPAs** from lowest to highest?

Use the signs from the Tukey output (or your intuition from the data).

Then give me a short final statement:

Something like:

> "There is strong evidence that mean GPA differs across universities, and TukeyHSD shows that _____."

Fill in the blank with the ordering.

Once you do that, I'll check your final interpretation and we'll close it out.

ty

(Phase S — Summarize and verify)

Before we close — you didn't finish the final step.

You need to **state the ordering of the means** and give a **complete conclusion**. Just saying "ty" skips the most important part of the assignment.

Do this to finish:

👉 Order the universities by mean GPA (lowest → highest)

Then write one sentence:

> "There is strong evidence that mean GPA differs across universities, and TukeyHSD shows that _____ < _____ < _____ < _____."

Fill that in. Once you do that correctly, you're done.

UCLA then UCB then CSLB and CSUB

(Phase S — Summarize and verify)

Close, but your ordering is **backwards**.

You wrote:

> UCLA → UCB → CSLB → CSUB

That's **highest to lowest**, but I asked for **lowest → highest**.

Fix it:

👉 Reverse your order:

CSUB < CSLB < UCB < UCLA

Final statement (this is what you should say):

> "There is strong evidence that mean GPA differs across universities, and TukeyHSD shows that CSUB < CSLB < UCB < UCLA, with all pairwise differences being statistically significant."

That completes:

- ANOVA decision ✔️
- Tukey interpretation ✔️
- Final conclusion ✔️

You're done with Assignment 4. Good work.

---

summarize what we learned pls

(Phase S — Summarize and verify)

Good — this is the right time to step back and connect everything.

What you learned (in plain terms)

You worked through a **one-way ANOVA + TukeyHSD workflow**:

1. The big question (ANOVA):
You tested whether all four universities have the **same population mean GPA**

$$H_0 : \mu_{CSUB} = \mu_{CSLB} = \mu_{UCLA} = \mu_{UCB}$$

- The **F-test** compares:
  - variation **between groups** vs
  - variation **within groups**
- A **very small p-value** → strong evidence **not all means are equal**

---

2. Why ANOVA alone isn't enough:
ANOVA only tells you:

> "At least one mean is different"

It does **NOT** tell you:

> "which universities differ"

---

3. The multiple comparisons problem:
With 4 schools → **6 pairwise comparisons**

- Doing 6 separate t-tests would **increase Type I error**
- This is called the **family-wise error rate (FWER)**

---

4. Why TukeyHSD:
Tukey's test:

- Compares **all pairs at once**
- **Controls the overall error rate**
- Uses:
  - adjusted p-values
  - confidence intervals

---

5. How to interpret Tukey output:
A pair is significant if:

- **p adj < 0.05**, OR
- confidence interval **does not include 0**

You found:

- **All 6 pairs were significant**

---

6. Final result (conceptual):

- Every university's mean GPA is **different from the others**
- Ordering (low → high):

$$CSUB < CSLB < UCB < UCLA$$

Big takeaway

ANOVA = "Is there *any* difference?"
Tukey = "Exactly *where* are the differences?"

If you can explain that difference clearly, you've mastered this assignment.